

Solution — Planks

1 Modelling

The problem states that you are given (one-dimensional) planks of different lengths and you are asked in how many different ways you can create a square by using all of the planks. Let x_1, \dots, x_n denote the lengths of the n planks. Since the arrangement of the planks within one side of the square does not matter, each of the four sides can be represented by a subset $S_i \subseteq \{x_1, \dots, x_n\}$, for $i \in \{1, \dots, 4\}$. Moreover, as one plank can not be used twice and at the same time all the planks must be used we have

$$S_1 \cup S_2 \cup S_3 \cup S_4 = \{x_1, \dots, x_n\}, \quad (1)$$

for all $1 \leq i < j \leq 4$,

$$S_i \cap S_j = \emptyset, \quad (2)$$

and for each $k \in \{1, 2, 3, 4\}$

$$\sum_{x' \in S_k} x' = S/4, \quad (3)$$

where $S := x_1 + \dots + x_n$. The last equation comes from the fact that the total sum of the plank lengths within a side must be equal for each side, as they together form a square. Thus, the number of ways we can partition the set $\{x_1, \dots, x_n\}$ into four pairwise disjoint sets $\{x_1, \dots, x_n\} = S_1 \cup S_2 \cup S_3 \cup S_4$ such that (3) holds, denoted by X , is almost what the problem asks from you. The last detail that one needs to notice is that for a given 4-tuple of sets (S_1, S_2, S_3, S_4) that satisfy (1), (2), and (3), any relabelling of the set indices produces the same square. As there are $4! = 24$ different ways to label a 4-tuple, we conclude that the correct output is $X/24$.

2 Algorithm Design

By looking at the constraints we see that n is at most 20. This suggests that a solution running in exponential time might be sufficient. Let us start by analysing a brute-force approach.

$O(n \cdot 4^n)$ solution (20 points). In how many ways can we partition a set of n elements into four disjoint subsets? Since each element is a member of one of the four sets, there are exactly 4^n different partitions. Since not all of the partitions satisfy (3) (they satisfy (1) and (2) by definition), we need to discard those that do not. Generating the partitions of $\{x_1, \dots, x_n\}$ consisting of 4 sets and checking which partition satisfies (3) can be done in $O(n \cdot 4^n)$ time. For $n = 20$ such algorithm would time out, but we can attack the first test set which guarantees $n \leq 8$. Furthermore, since the first two test sets also guarantee that the answer is always 0 or 1, we can output 1 as soon as we find one partition with the desired properties or output 0 if we find no such partition.

$O(n \cdot 2^n)$ solution (60–100 points). Using the *Split and List* technique presented in the tutorial of week 5, one can achieve a much faster running time than in the previous solution. The Split and List technique was illustrated on the problem Subset Sum. In that problem you are given a set X and a number k and your task is to figure out if there exists a subset $S \subseteq X$ such that the sum of the elements in S is equal to k . This problem can be reformulated in a slightly different way: is there a partition of X into two disjoint sets $X = S_1 \cup S_2$ such that

$$\sum_{x' \in S_1} x' = k \quad \text{and} \quad \sum_{x'' \in S_2} x'' = S - k,$$

where S is the total sum of elements in X . This formulation looks quite similar to our current problem and we can try to use the same strategy as for solving Subset Sum.

Let $X = X_1 \cup X_2$ be an arbitrary partition of set X into two disjoint sets such that $|X_1| = |X_2| = n/2$ (for simplicity we assume n is even). Let $i \in \{1, 2\}$ and let F_i be defined as follows

$$F_i = \{(s_1, s_2, s_3, s_4) : \exists \text{ partition of } X_i \text{ into four disjoint sets } X_i = P_1 \cup P_2 \cup P_3 \cup P_4 \\ \text{such that for all } j \in \{1, 2, 3, 4\} \text{ we have } \sum_{x' \in P_j} x' = s_j\}. \quad (4)$$

Having F_1 and F_2 it is not hard to see that the following algorithm correctly determines if our problem has 0 or more solutions (but not the exact number)

```
1 int result = 0
2 SQ = (x1 + ... + xn) / 4
3 for every 4-tuple (s1, s2, s3, s4) in F_1:
4     if there exist (SQ - s1, SQ - s2, SQ - s3, SQ - s4) in F_2:
5         result = 1
```

Thus if we can test a membership of a 4-tuple in F_2 in time t , then the previous algorithm can be implemented in time roughly equal to $t \cdot |F_1|$. By using an appropriate data structure, e.g. set, testing for membership can be done in time $O(\log |F_2|)$. The size of set F_1 is at most $4^{n/2}$ as there are at most $4^{n/2}$ partitions of X_1 into four sets and thus the previous algorithm runs in time

$$O(4^{n/2} \log 4^{n/2}) = O(n \cdot 2^n).$$

This approach achieves 60 points.

In order to get the exact number of solutions (and consequently full points), we need to modify our algorithm a bit. First, we need that F_2 is stored using a data structure which supports multiple copies of the same element, i.e. a vector, map, etc. The next algorithm computes the number of solutions

```
1 int result = 0
2 SQ = (x1 + ... + xn) / 4
3 for every 4-tuple (s1, s2, s3, s4) in F_1:
4     counter = number of copies of (SQ - s1, SQ - s2, SQ - s3, SQ - s4) in F_2
5     result += counter
6
7 result /= 24
```

Although data structure `std::multiset` sounds like a natural candidate, we advise *against* using it. Counting the number of copies of an element in `std::multiset` runs in linear time in the number of copies of the element. As we would like to execute line 4 in time $\log |F_2|$ this is

not good enough. Instead, we can store F_2 in a vector and sort it before iterating through F_1 . Now, we can execute line 4 in time $O(\log|F_2|)$ by using binary search.

Remark: Although the approach described above correctly implemented scores 100 points, there is a way to easily improve the performance of the algorithm by a factor of roughly two. Because of the symmetry of the problem, we can fill F_1 only with 4-tuples such that x_1 is always in P_1 . This cuts the size of F_1 in four. In this way we remove some of the symmetries, so in the end you have to divide result by 6 instead of 24.

3 Implementation

Let us first discuss a possible way to implement the brute-force solution. Probably the easiest solution is to use a **back-track approach**. We can use `std::vector<std::vector<int> > assignment` to store which elements are assigned to which set, where `assignment[i]` contains all elements which are part of S_i . Having this, the following recursive function returns the number of solutions.

```

1 int back_track(int id) {
2     if (id >= n) {
3         bool ok = true;
4         for (int i = 0; i < 4; ++i) {
5             int sum = 0;
6             for (int j = 0; j < (int)assignment[i].size(); ++j)
7                 sum += planks[assignment[i][j]];
8
9             if (sum != side_length) {
10                ok = false;
11                break;
12            }
13        }
14        if (ok) return 1;
15        else return 0;
16    }
17
18    int counter = 0;
19    for (int side = 0; side < 4; ++side) {
20        assignment[side].push_back(id);
21        counter += back_track(id + 1);
22        assignment[side].pop_back();
23    }
24
25    return counter;
26 }
```

Let us now see how to approach the implementation of $O(n \cdot 2^n)$ solution. We can represent F_1, F_2 as `std::vector<std::vector<int> > F1, F2`. Both of the vectors can be filled with a similar recursive procedure as in the back-track solution above. After we fill both vectors we need to sort F_2 which can be done as simple as

```
1 std::sort(F2.begin(), F2.end());
```

Finally, we want to check how many copies of a particular 4-tuple is contained in F_2 . For this we can use functions `std::equal_range` (which runs in $O(\log|F_2|)$) and `std::distance` as follows.

```

1 typedef std::vector<int> VI;
2 typedef std::vector<VI> VVI;
3
4 // Variable target contains our target 4-tuple
5 std::pair<VVI::iterator, VVI::iterator> bounds;
6 bounds = std::equal_range(F2.begin(), F2.end(), target);
7 long long counter = std::distance(bounds.first, bounds.second);

```

4 Appendix

The following code is an implementation of $O(n \cdot 2^n)$ solution explained above.

```

1 #include <iostream>
2 #include <vector>
3 #include <algorithm>
4
5 typedef std::vector<int> VI;
6 typedef std::vector<VI> VVI;
7
8 void back_track(int id, int ubound, VVI &F, VVI &assignment, const VI &planks) {
9     if (id >= ubound) {
10         VI tuple(4, 0);
11         for (int i = 0; i < 4; ++i) {
12             for (int j = 0; j < (int)assignment[i].size(); ++j)
13                 tuple[i] += planks[assignment[i][j]];
14         }
15         F.push_back(tuple);
16         return;
17     }
18
19     for (int i = 0; i < 4; ++i) {
20         assignment[i].push_back(id); // Try to put id-th plank to i-th set
21         back_track(id + 1, ubound, F, assignment, planks); // Recurse
22         assignment[i].pop_back(); // Remove id-th plank from i-th set
23     }
24 }
25
26 void testcase() {
27     int n; std::cin >> n;
28     std::vector<int> planks;
29     for (int i = 0; i < n; ++i) {
30         int plank; std::cin >> plank;
31         planks.push_back(plank);
32     }
33
34     int sum = 0;
35     long long result = 0;
36
37     for (int i = 0; i < (int)planks.size(); ++i)
38         sum += planks[i];
39
40     // If the total sum of lengths is not divisible by four, it is not possible
41     // to create a square.
42     if (sum % 4 != 0) {
43         std::cout << 0 << std::endl;
44         return;
45     }

```

```

46
47 VVI F1, assignment(4);
48 // Generate all 4-tuple for the first half of the set.
49 back_track(0, n/2, F1, assignment, planks);
50
51 VVI F2, assignment2(4);
52 // Generate all 4-tuple for the second half of the set.
53 back_track(n/2, n, F2, assignment2, planks);
54 std::sort(F2.begin(), F2.end());
55
56 for (int idx = 0; idx < (int)F1.size(); ++idx) {
57     std::vector<int> member = F1[idx];
58     for (int i = 0; i < 4; ++i)
59         member[i] = sum/4 - member[i];
60
61     std::pair<VVI::iterator, VVI::iterator> bounds;
62     bounds = std::equal_range(F2.begin(), F2.end(), member);
63     long long counter = std::distance(bounds.first, bounds.second);
64     result += counter;
65 }
66 std::cout << result / 24 << std::endl;
67 }
68
69 int main() {
70     std::ios_base::sync_with_stdio(false);
71
72     int t; std::cin >> t;
73     for (int i = 0; i < t; i++) {
74         testcase();
75     }
76 }

```